

ML-1 Graded Assessment

DATA_SET:

A Chinese automobile company Geely Auto contracted an automobile consulting company to understand the factors on which the pricing of cars depends. Specifically, they want to understand the factors affecting the pricing of cars in the American market, since those may be very different from the Chinese market. The company wants to know

Based on various market surveys, the consulting firm has gathered a large dataset of different types of cars across the American market.

- Car_ID== Unique id of each observation (Integer)
- Symboling== Its assigned insurance risk rating, A value of +3 indicates that the auto is risky, -3 that it is probably pretty safe.(Categorical)
- carCompany== Name of car company (Categorical)
- fueltype== Car fuel type i.e gas or diesel (Categorical)
- aspiration== Aspiration used in a car (Categorical)
- doornumber== Number of doors in a car (Categorical)
- carbody== body of car (Categorical)
- drivewheel== type of drive wheel (Categorical)
- enginelocation== Location of car engine (Categorical)
- wheelbase== Wheelbase of car (Numeric)
- carlength== Length of car (Numeric)
- carwidth== Width of car (Numeric)
- carheight== height of car (Numeric)
- curbweight== The weight of a car without occupants or baggage. (Numeric)
- enginetype== Type of engine. (Categorical)
- cylindernumber== cylinder placed in the car (Categorical)
- enginesize== Size of car (Numeric)
- fuelsystem== Fuel system of car (Categorical)
- boreratio== Boreratio of car (Numeric)
- stroke== Stroke or volume inside the engine (Numeric)
- compressionratio== compression ratio of car (Numeric)

- horsepower== Horsepower (Numeric)
- peakrpm== car peak rpm (Numeric)
- citympg== Mileage in city (Numeric)
- highwaympg== Mileage on highway (Numeric)
- price(Dependent variable)== Price of car (Numeric)

```
In [27]: import warnings
warnings.filterwarnings('ignore')
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

1. Data Understanding (5 marks)

a. Read the dataset (tab, csv, xls, txt, inbuilt dataset). What are the number of rows and no. of cols & types of variables (continuous, categorical etc.)? (1 MARK)

b. Calculate five-point summary for numerical variables (1 MARK)

c. Summarize observations for categorical variables – no. of categories, % observations in each category. (1 mark)

d. Check for defects in the data such as missing values, null, outliers, etc. (2 marks)

a. Read the dataset (tab, csv, xls, txt, inbuilt dataset). What are the number of rows and no. of cols & types of variables (continuous, categorical etc.)? (1 MARK)

```
In [29]: cars = pd.read_csv('Car_Data.csv')
cars.head()
```

```
Out[29]:
```

| | car_ID | symboling | CarName | fueltype | aspiration | doornumber | carbody | drivewheel | enginelocation | wheelbase | ... | engine size | fuelsystem | bore |
|---|--------|-----------|-----------------------------|----------|------------|------------|-------------|------------|----------------|-----------|-----|-------------|------------|------|
| 0 | 1 | 3 | alfa-romero giulia | gas | std | two | convertible | rwd | front | 88.6 | ... | 130 | mpfi | |
| 1 | 2 | 3 | alfa-romero stelvio | gas | std | two | convertible | rwd | front | 88.6 | ... | 130 | mpfi | |
| 2 | 3 | 1 | alfa-romero Quadrifoglio | gas | std | two | hatchback | rwd | front | 94.5 | ... | 152 | mpfi | |

| | car_ID | symboling | CarName | fueltype | aspiration | doornumber | carbody | drivewheel | enginelocation | wheelbase | ... | enginesize | fuelsystem | bore |
|--|--------|-----------|---------|-------------|------------|------------|---------|------------|----------------|-----------|----------|------------|------------|------|
| | 3 | 4 | 2 | audi 100 ls | gas | std | four | sedan | fwd | front | 99.8 ... | 109 | mpfi | |
| | 4 | 5 | 2 | audi 100ls | gas | std | four | sedan | 4wd | front | 99.4 ... | 136 | mpfi | |

5 rows × 26 columns



In [30]: `cars.shape`

Out[30]: (205, 26)

In [31]: `cat_cols = cars.select_dtypes(include = 'object')`
`num_cols = cars.select_dtypes(include = np.number)`
`print('Continuous variables are : ', num_cols.columns)`
`print('Categorical Columns are : ', cat_cols.columns)`

Continuous variables are : Index(['car_ID', 'symboling', 'wheelbase', 'carlength', 'carwidth',
'carheight', 'curbweight', 'enginesize', 'bore_ratio', 'stroke',
'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
'price'],
dtype='object')

Categorical Columns are : Index(['CarName', 'fueltype', 'aspiration', 'doornumber', 'carbody',
'drivewheel', 'engine_location', 'engine_type', 'cylindernumber',
'fuelsystem'],
dtype='object')

b. Calculate five-point summary for numerical variables (1 MARK)

In [32]: `cars.describe()`

| | car_ID | symboling | wheelbase | carlength | carwidth | carheight | curbweight | enginesize | bore_ratio | stroke | compressionratio | horsepower |
|--------------|------------|------------|------------|------------|------------|------------|-------------|------------|------------|------------|------------------|------------|
| count | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 |
| mean | 103.000000 | 0.834146 | 98.756585 | 174.049268 | 65.907805 | 53.724878 | 2555.565854 | 126.907317 | 3.329756 | 3.255415 | 10.142537 | 104.111111 |
| std | 59.322565 | 1.245307 | 6.021776 | 12.337289 | 2.145204 | 2.443522 | 520.680204 | 41.642693 | 0.270844 | 0.313597 | 3.972040 | 39.544444 |
| min | 1.000000 | -2.000000 | 86.600000 | 141.100000 | 60.300000 | 47.800000 | 1488.000000 | 61.000000 | 2.540000 | 2.070000 | 7.000000 | 48.000000 |
| 25% | 52.000000 | 0.000000 | 94.500000 | 166.300000 | 64.100000 | 52.000000 | 2145.000000 | 97.000000 | 3.150000 | 3.110000 | 8.600000 | 70.000000 |

| | car_ID | symboling | wheelbase | carlength | carwidth | carheight | curbweight | enginesize | boreratio | stroke | compressionratio | horsepower |
|------------|------------|-----------|------------|------------|-----------|-----------|-------------|------------|-----------|----------|------------------|------------|
| 50% | 103.000000 | 1.000000 | 97.000000 | 173.200000 | 65.500000 | 54.100000 | 2414.000000 | 120.000000 | 3.310000 | 3.290000 | 9.000000 | 95.000000 |
| 75% | 154.000000 | 2.000000 | 102.400000 | 183.100000 | 66.900000 | 55.500000 | 2935.000000 | 141.000000 | 3.580000 | 3.410000 | 9.400000 | 116.000000 |
| max | 205.000000 | 3.000000 | 120.900000 | 208.100000 | 72.300000 | 59.800000 | 4066.000000 | 326.000000 | 3.940000 | 4.170000 | 23.000000 | 288.000000 |

c. Summarize observations for categorical variables – no. of categories, % observations in each category. (1 mark)

In [33]: `cars.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling              205 non-null   int64
2   CarName               205 non-null   object
3   fueltype              205 non-null   object
4   aspiration             205 non-null   object
5   doornumber            205 non-null   object
6   carbody               205 non-null   object
7   drivewheel            205 non-null   object
8   enginelocation        205 non-null   object
9   wheelbase             205 non-null   float64
10  carlength             205 non-null   float64
11  carwidth              205 non-null   float64
12  carheight             205 non-null   float64
13  curbweight            205 non-null   int64
14  enginetype            205 non-null   object
15  cylindernumber        205 non-null   object
16  enginesize            205 non-null   int64
17  fuelsystem            205 non-null   object
18  boreratio             205 non-null   float64
19  stroke                205 non-null   float64
20  compressionratio      205 non-null   float64
21  horsepower            205 non-null   int64
22  peakrpm              205 non-null   int64
23  citympg               205 non-null   int64
24  highwaympg           205 non-null   int64
25  price                205 non-null   float64
```

```
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

```
In [34]: print('% Values in each categorical columns')
        for i in cat_cols.columns:
            print('\n% Values in column ',i)
            print((cars[i].value_counts()/len(cars[i])*100))
```

```
% Values in each categorical columns
```

```
% Values in column  CarName
```

```
peugeot 504          2.926829
toyota corolla       2.926829
toyota corona        2.926829
subaru dl            1.951220
mitsubishi outlander 1.463415
```

```
...
```

```
buick skyhawk        0.487805
subaru tribeca        0.487805
buick electra 225 custom 0.487805
alfa-romero stelvio  0.487805
volkswagen rabbit    0.487805
```

```
Name: CarName, Length: 147, dtype: float64
```

```
% Values in column  fueltype
```

```
gas      90.243902
diesel   9.756098
Name: fueltype, dtype: float64
```

```
% Values in column  aspiration
```

```
std      81.95122
turbo    18.04878
Name: aspiration, dtype: float64
```

```
% Values in column  doornumber
```

```
four     56.097561
two      43.902439
Name: doornumber, dtype: float64
```

```
% Values in column  carbody
```

```
sedan     46.829268
hatchback 34.146341
wagon     12.195122
hardtop    3.902439
convertible 2.926829
Name: carbody, dtype: float64
```

```
% Values in column drivewheel
```

```
fwd    58.536585
```

```
rwd    37.073171
```

```
4wd     4.390244
```

```
Name: drivewheel, dtype: float64
```

```
% Values in column enginelocation
```

```
front   98.536585
```

```
rear     1.463415
```

```
Name: enginelocation, dtype: float64
```

```
% Values in column enginetype
```

```
ohc     72.195122
```

```
ohcf     7.317073
```

```
ohcv     6.341463
```

```
dohc     5.853659
```

```
l         5.853659
```

```
rotor     1.951220
```

```
dohcv     0.487805
```

```
Name: enginetype, dtype: float64
```

```
% Values in column cylindernumber
```

```
four    77.560976
```

```
six     11.707317
```

```
five     5.365854
```

```
eight    2.439024
```

```
two      1.951220
```

```
twelve   0.487805
```

```
three    0.487805
```

```
Name: cylindernumber, dtype: float64
```

```
% Values in column fuelsystem
```

```
mpfi    45.853659
```

```
2bbl    32.195122
```

```
idi     9.756098
```

```
1bbl    5.365854
```

```
spdi    4.390244
```

```
4bbl    1.463415
```

```
mfi     0.487805
```

```
spfi    0.487805
```

```
Name: fuelsystem, dtype: float64
```

```
In [35]: print('% Values in each categorical columns')
          for i in cat_cols.columns:
              print('\n% Values in column ',i)
              print((cars[i].value_counts(1)))
```

% Values in each categorical columns

% Values in column CarName

peugeot 504 0.029268

toyota corolla 0.029268

toyota corona 0.029268

subaru dl 0.019512

mitsubishi outlander 0.014634

...

buick skyhawk 0.004878

subaru tribeca 0.004878

buick electra 225 custom 0.004878

alfa-romero stelvio 0.004878

volkswagen rabbit 0.004878

Name: CarName, Length: 147, dtype: float64

% Values in column fueltype

gas 0.902439

diesel 0.097561

Name: fueltype, dtype: float64

% Values in column aspiration

std 0.819512

turbo 0.180488

Name: aspiration, dtype: float64

% Values in column doornumber

four 0.560976

two 0.439024

Name: doornumber, dtype: float64

% Values in column carbody

sedan 0.468293

hatchback 0.341463

wagon 0.121951

hardtop 0.039024

convertible 0.029268

Name: carbody, dtype: float64

% Values in column drivewheel

fwd 0.585366

rwd 0.370732

4wd 0.043902

Name: drivewheel, dtype: float64

% Values in column enginelocation

front 0.985366

```
rear      0.014634
Name: enginelocation, dtype: float64
```

```
% Values in column  enginetype
```

```
ohc      0.721951
ohcf     0.073171
ohcv     0.063415
dohc     0.058537
l        0.058537
rotor    0.019512
dohcv    0.004878
```

```
Name: enginetype, dtype: float64
```

```
% Values in column  cylindernumber
```

```
four     0.775610
six      0.117073
five     0.053659
eight    0.024390
two      0.019512
twelve   0.004878
three    0.004878
```

```
Name: cylindernumber, dtype: float64
```

```
% Values in column  fuelsystem
```

```
mpfi     0.458537
2bbl     0.321951
idi      0.097561
1bbl     0.053659
spdi     0.043902
4bbl     0.014634
mfi      0.004878
spfi     0.004878
```

```
Name: fuelsystem, dtype: float64
```

d. Check for defects in the data such as missing values, null, outliers, etc. (2 marks)

```
In [36]: cars.isnull().sum()
```

```
Out[36]: car_ID      0
symboling    0
CarName      0
fueltype     0
aspiration   0
doornumber   0
carbody      0
drivewheel   0
enginelocation 0
```



```

wheelbase      0
carlength      0
carwidth       0
carheight      0
curbweight     0
enginetype     0
cylindernumber 0
enginesize     0
fuelsystem     0
boretostroke   0
stroke         0
compressionratio 0
horsepower     0
peakrpm        0
citympg        0
highwaympg     0
price          0
dtype: int64

```


- As we can see, this dataset does not contain any missing value


* Outliers:


```

In [37]: plt.figure(figsize=(10,20))
for i,col in enumerate(num_cols,1):
    plt.subplot(16,1,i)
    sns.boxplot(cars[col])
    plt.ylabel(col)
plt.show()

```







Data Cleaning

```
In [38]: cars['CarName'].unique()
```

```
Out[38]: array(['alfa-romero giulia', 'alfa-romero stelvio',
                'alfa-romero Quadrifoglio', 'audi 100 ls', 'audi 100ls',
                'audi fox', 'audi 5000', 'audi 4000', 'audi 5000s (diesel)',
                'bmw 320i', 'bmw x1', 'bmw x3', 'bmw z4', 'bmw x4', 'bmw x5',
                'chevrolet impala', 'chevrolet monte carlo', 'chevrolet vega 2300',
                'dodge rampage', 'dodge challenger se', 'dodge d200',
                'dodge monaco (sw)', 'dodge colt hardtop', 'dodge colt (sw)',
                'dodge coronet custom', 'dodge dart custom',
                'dodge coronet custom (sw)', 'honda civic', 'honda civic cvcc',
                'honda accord cvcc', 'honda accord lx', 'honda civic 1500 gl',
                'honda accord', 'honda civic 1300', 'honda prelude',
                'honda civic (auto)', 'isuzu MU-X', 'isuzu D-Max',
                'isuzu D-Max V-Cross', 'jaguar xj', 'jaguar xf', 'jaguar xk',
                'maxda rx3', 'maxda glc deluxe', 'mazda rx2 coupe', 'mazda rx-4',
                'mazda glc deluxe', 'mazda 626', 'mazda glc', 'mazda rx-7 gs',
                'mazda glc 4', 'mazda glc custom l', 'mazda glc custom',
                'buick electra 225 custom', 'buick century luxury (sw)',
                'buick century', 'buick skyhawk', 'buick opel isuzu deluxe',
                'buick skylark', 'buick century special',
                'buick regal sport coupe (turbo)', 'mercury cougar',
                'mitsubishi mirage', 'mitsubishi lancer', 'mitsubishi outlander',
                'mitsubishi g4', 'mitsubishi mirage g4', 'mitsubishi montero',
                'mitsubishi pajero', 'Nissan versa', 'nissan gt-r', 'nissan rogue',
                'nissan latio', 'nissan titan', 'nissan leaf', 'nissan juke',
                'nissan note', 'nissan clipper', 'nissan nv200', 'nissan dayz',
                'nissan fuga', 'nissan otti', 'nissan teana', 'nissan kicks',
                'peugeot 504', 'peugeot 304', 'peugeot 504 (sw)', 'peugeot 604sl',
                'peugeot 505s turbo diesel', 'plymouth fury iii',
                'plymouth cricket', 'plymouth satellite custom (sw)',
                'plymouth fury gran sedan', 'plymouth valiant', 'plymouth duster',
```

```
'porsche macan', 'porcshce panamera', 'porsche cayenne',
'porsche boxter', 'renault 12tl', 'renault 5 gtl', 'saab 99e',
'saab 99le', 'saab 99gle', 'subaru', 'subaru dl', 'subaru brz',
'subaru baja', 'subaru r1', 'subaru r2', 'subaru trezia',
'subaru tribeca', 'toyota corona mark ii', 'toyota corona',
'toyota corolla 1200', 'toyota corona hardtop',
'toyota corolla 1600 (sw)', 'toyota carina', 'toyota mark ii',
'toyota corolla', 'toyota corolla liftback',
'toyota celica gt liftback', 'toyota corolla tercel',
'toyota corona liftback', 'toyota starlet', 'toyota tercel',
'toyota cressida', 'toyota celica gt', 'toyouta tercel',
'vokswagen rabbit', 'volkswagen 1131 deluxe sedan',
'volkswagen model 111', 'volkswagen type 3', 'volkswagen 411 (sw)',
'volkswagen super beetle', 'volkswagen dasher', 'vw dasher',
'vw rabbit', 'volkswagen rabbit', 'volkswagen rabbit custom',
'volvo 145e (sw)', 'volvo 144ea', 'volvo 244dl', 'volvo 245',
'volvo 264gl', 'volvo diesel', 'volvo 246'], dtype=object)
```

```
In [39]: #Splitting company name from CarName column
CompanyName = cars['CarName'].apply(lambda x : x.split(' ')[0])
cars.insert(3, "CompanyName", CompanyName)
cars.drop(['CarName'], axis=1, inplace=True)
cars.head()
```

```
Out[39]:
```

| | car_ID | symboling | CompanyName | fueltype | aspiration | doornumber | carbody | drivewheel | engine location | wheelbase | ... | enginesize | fuelsystem | b |
|---|--------|-----------|-------------|----------|------------|------------|-------------|------------|-----------------|-----------|-----|------------|------------|---|
| 0 | 1 | 3 | alfa-romero | gas | std | two | convertible | rwd | front | 88.6 | ... | 130 | mpfi | |
| 1 | 2 | 3 | alfa-romero | gas | std | two | convertible | rwd | front | 88.6 | ... | 130 | mpfi | |
| 2 | 3 | 1 | alfa-romero | gas | std | two | hatchback | rwd | front | 94.5 | ... | 152 | mpfi | |
| 3 | 4 | 2 | audi | gas | std | four | sedan | fwd | front | 99.8 | ... | 109 | mpfi | |
| 4 | 5 | 2 | audi | gas | std | four | sedan | 4wd | front | 99.4 | ... | 136 | mpfi | |

5 rows × 26 columns



```
In [40]: cars.CompanyName.unique()
```

```
Out[40]: array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',
               'isuzu', 'jaguar', 'maxda', 'mazda', 'buick', 'mercury',
               'mitsubishi', 'Nissan', 'nissan', 'peugeot', 'plymouth', 'porsche',
```

```
'porcshce', 'renault', 'saab', 'subaru', 'toyota', 'toyouta',
'volkswagen', 'volkswagen', 'vw', 'volvo'], dtype=object)
```

Fixing invalid values

- There seems to be some spelling error in the CompanyName column.
 - maxda = mazda
 - Nissan = nissan
 - porsche = porcshce
 - toyota = toyouta
 - volkswagen = volkswagen = vw

```
In [41]: cars.CompanyName = cars.CompanyName.str.lower()

def replace_name(a,b):
    cars.CompanyName.replace(a,b,inplace=True)

replace_name('maxda','mazda')
replace_name('porcshce','porsche')
replace_name('toyouta','toyota')
replace_name('volkswagen','volkswagen')
replace_name('vw','volkswagen')

cars.CompanyName.unique()
```



```
Out[41]: array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',
               'isuzu', 'jaguar', 'mazda', 'buick', 'mercury', 'mitsubishi',
               'nissan', 'peugeot', 'plymouth', 'porsche', 'renault', 'saab',
               'subaru', 'toyota', 'volkswagen', 'volvo'], dtype=object)
```

```
In [42]: plt.figure(figsize=(20,8))

plt.subplot(1,2,1)
plt.title('Car Price Distribution Plot')
sns.distplot(cars.price)

plt.subplot(1,2,2)
plt.title('Car Price Spread')
sns.boxplot(y=cars.price)

plt.show()
```



```
In [43]: print(cars.price.describe(percentiles = [0.25,0.50,0.75,0.85,0.90,1]))
```

```
count      205.000000
mean       13276.710571
std         7988.852332
min         5118.000000
25%         7788.000000
50%        10295.000000
75%        16503.000000
85%        18500.000000
90%        22563.000000
100%       45400.000000
max        45400.000000
Name: price, dtype: float64
```

Inference :

1. The plot seemed to be right-skewed, meaning that the most prices in the dataset are low(Below 15,000).
2. There is a significant difference between the mean and the median of the price distribution.

3. The data points are far spread out from the mean, which indicates a high variance in the car prices.(85% of the prices are below 18,500, whereas the remaining 15% are between 18,500 and 45,400.)

2. Data Preparation (15 marks)

- Fix the defects found above and do appropriate treatment if any. (5 marks)
- Visualize the data using relevant plots. Find out the variables which are highly correlated with target variable? (5 marks)
- Do you want to exclude some variables from the model based on this analysis? What other actions will you take? (2 marks)
- Split dataset into train and test (70:30). Are both train and test representative of the overall data? How would you ascertain this statistically? (3 marks)

```
In [44]: plt.figure(figsize=(25, 6))

plt.subplot(1,3,1)
plt1 = cars.CompanyName.value_counts().plot('bar')
plt.title('Companies Histogram')
plt1.set(xlabel = 'Car company', ylabel='Frequency of company')

plt.subplot(1,3,2)
plt1 = cars.fueltype.value_counts().plot('bar')
plt.title('Fuel Type Histogram')
plt1.set(xlabel = 'Fuel Type', ylabel='Frequency of fuel type')

plt.subplot(1,3,3)
plt1 = cars.carbody.value_counts().plot('bar')
plt.title('Car Type Histogram')
plt1.set(xlabel = 'Car Type', ylabel='Frequency of Car type')

plt.show()
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-44-63515b9191a3> in <module>
      2
      3 plt.subplot(1,3,1)
----> 4 plt1 = cars.CompanyName.value_counts().plot('bar')
      5 plt.title('Companies Histogram')
      6 plt1.set(xlabel = 'Car company', ylabel='Frequency of company')
```


```

~\anaconda3\lib\site-packages\pandas\plotting\_core.py in __call__(self, *args, **kwargs)
    875     plot_backend = _get_plot_backend(kwargs.pop("backend", None))
    876
--> 877     x, y, kind, kwargs = self._get_call_args(
    878         plot_backend.__name__, self._parent, args, kwargs
    879     )

~\anaconda3\lib\site-packages\pandas\plotting\_core.py in _get_call_args(backend_name, data, args, kwargs)
    859     f"`Series.plot({positional_args})`."
    860     )
--> 861     raise TypeError(msg)
    862
    863     pos_args = {name: value for value, (name, _) in zip(args, arg_def)}

```

TypeError: `Series.plot()` should not be called with positional arguments, only keyword arguments. The order of positional arguments will change in the future. Use `Series.plot(kind='bar')` instead of `Series.plot('bar',)`.



Inference :

1. Toyota seemed to be favored car company.
2. Number of gas fueled cars are more than diesel .

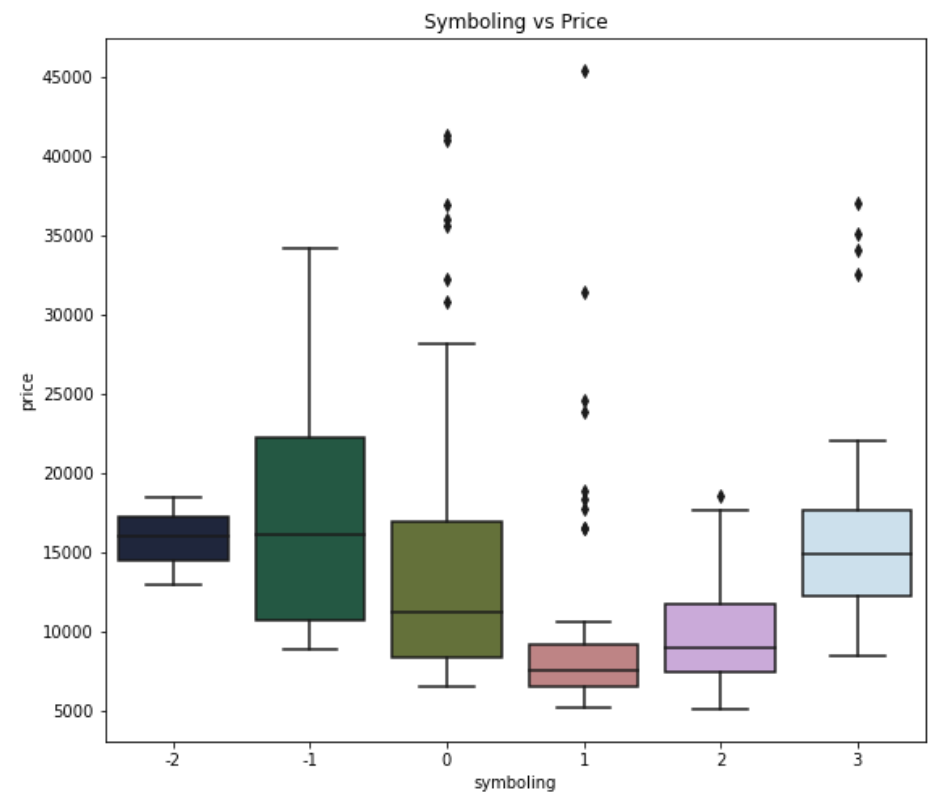

3. sedan is the top car type preferred.

```
In [45]: plt.figure(figsize=(20,8))

plt.subplot(1,2,1)
plt.title('Symboling Histogram')
sns.countplot(cars.symboling, palette="cubehelix")

plt.subplot(1,2,2)
plt.title('Symboling vs Price')
sns.boxplot(x=cars.symboling, y=cars.price, palette="cubehelix")

plt.show()
```



Inference :

1. It seems that the symboling with 0 and 1 values have high number of rows (i.e. They are most sold.)

2. The cars with -1 symboling seems to be high priced (as it makes sense too, insurance risk rating -1 is quite good). But it seems that symboling with 3 value has the price range similar to -2 value. There is a dip in price at symboling 1 .



```
In [46]: plt.figure(figsize=(20,8))


plt.subplot(1,2,1)
plt.title('Engine Type Histogram')
sns.countplot(cars.engine_type, palette="Blues_d")

plt.subplot(1,2,2)
plt.title('Engine Type vs Price')
sns.boxplot(x=cars.engine_type, y=cars.price, palette="PuBuGn")

plt.show()

df = pd.DataFrame(cars.groupby(['engine_type'])['price'].mean().sort_values(ascending = False))
df.plot.bar(figsize=(8,6))
plt.title('Engine Type vs Average Price')
plt.show()
```





Inference :

1. ohc Engine type seems to be most favored type.
2. ohcv has the highest price range (While dohc has only one row), ohc and ohcf have the low price range.


```
In [47]: plt.figure(figsize=(25, 6))



df = pd.DataFrame(cars.groupby(['CompanyName'])['price'].mean().sort_values(ascending = False))
df.plot.bar()
plt.title('Company Name vs Average Price')
plt.show()

df = pd.DataFrame(cars.groupby(['fueltype'])['price'].mean().sort_values(ascending = False))
df.plot.bar()
plt.title('Fuel Type vs Average Price')
plt.show()
```

```
df = pd.DataFrame(cars.groupby(['carbody'])['price'].mean().sort_values(ascending = False))  
df.plot.bar()  
plt.title('Car Type vs Average Price')  
plt.show()
```

<Figure size 1800x432 with 0 Axes>





Inference :

1. Jaguar and Buick seem to have highest average price.

2. diesel has higher average price than gas.
3. hardtop and convertible have higher average price.

```
In [48]: plt.figure(figsize=(15,5))

plt.subplot(1,2,1)
plt.title('Door Number Histogram')
sns.countplot(cars.doornumber, palette="plasma")

plt.subplot(1,2,2)
plt.title('Door Number vs Price')
sns.boxplot(x=cars.doornumber, y=cars.price, palette="plasma")

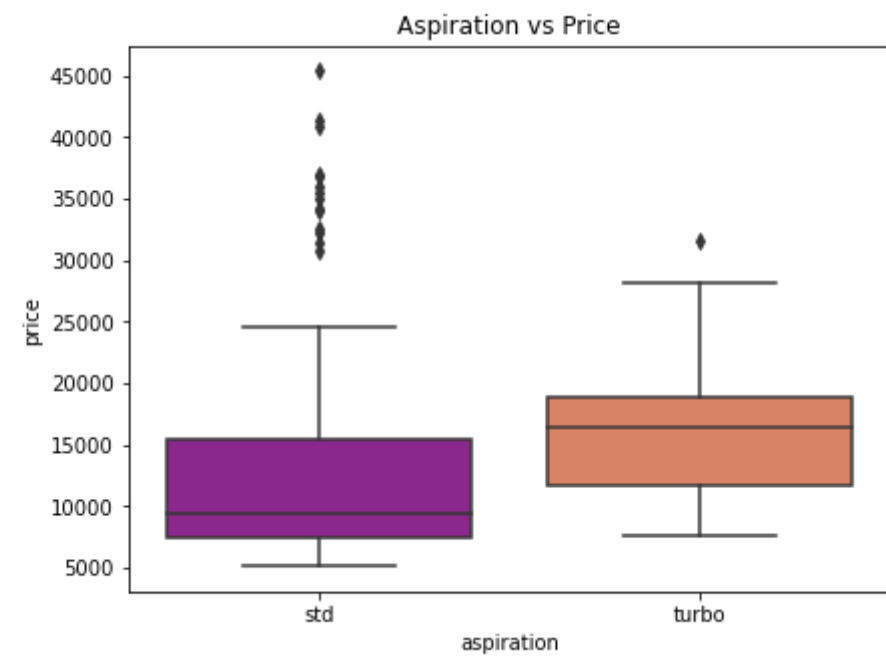


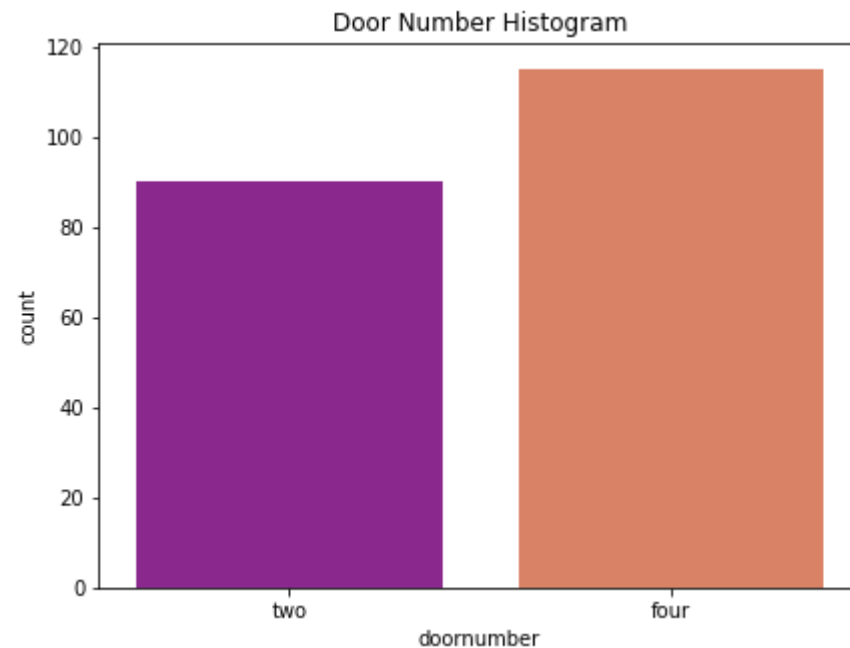
plt.show()

plt.figure(figsize=(15,5))

plt.subplot(1,2,1)
plt.title('Aspiration Histogram')
sns.countplot(cars.aspiration, palette="plasma")

plt.subplot(1,2,2)
plt.title('Aspiration vs Price')
sns.boxplot(x=cars.aspiration, y=cars.price, palette="plasma")

plt.show()
```



Inference :



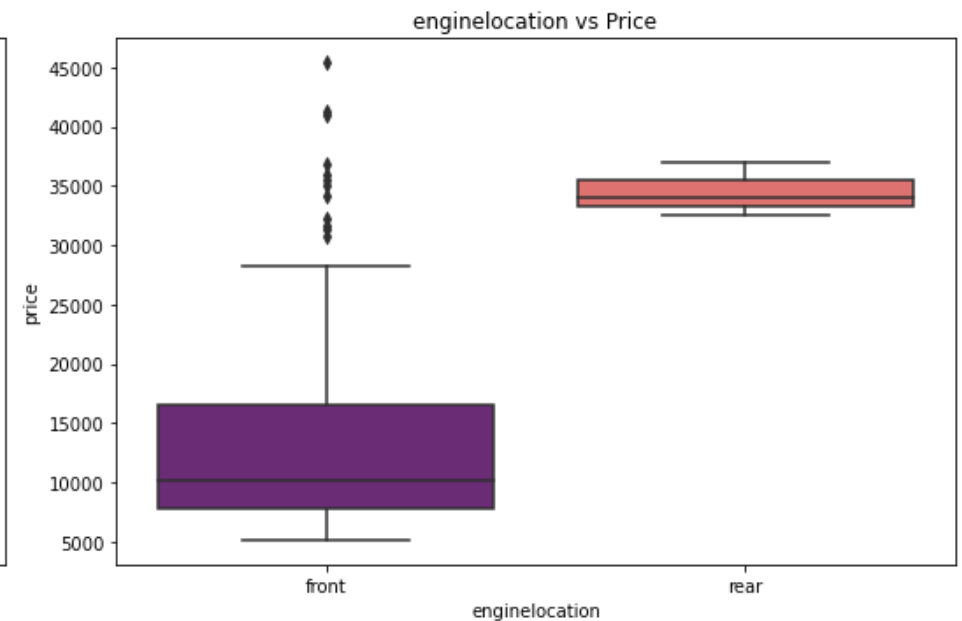
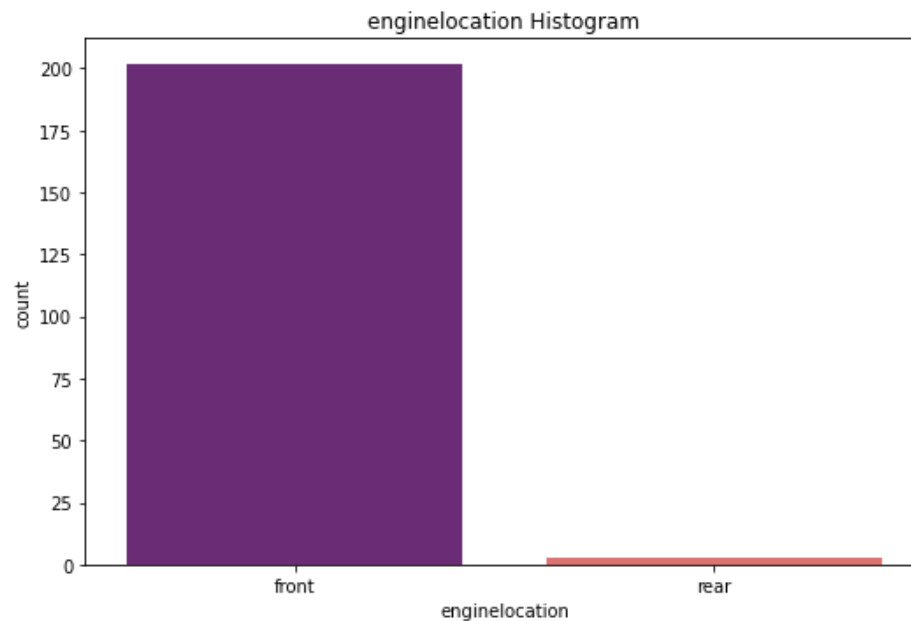
1. doornumber variable is not affecting the price much. There is no significant difference between the categories in it.
2. It seems aspiration with turbo have higher price range than the std (though it has some high values outside the whiskers.)



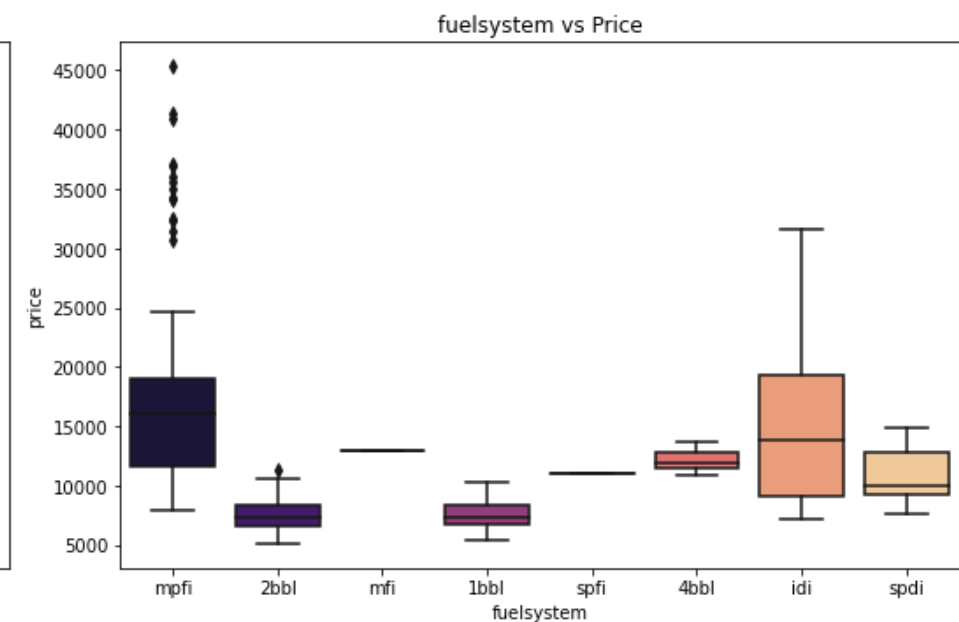
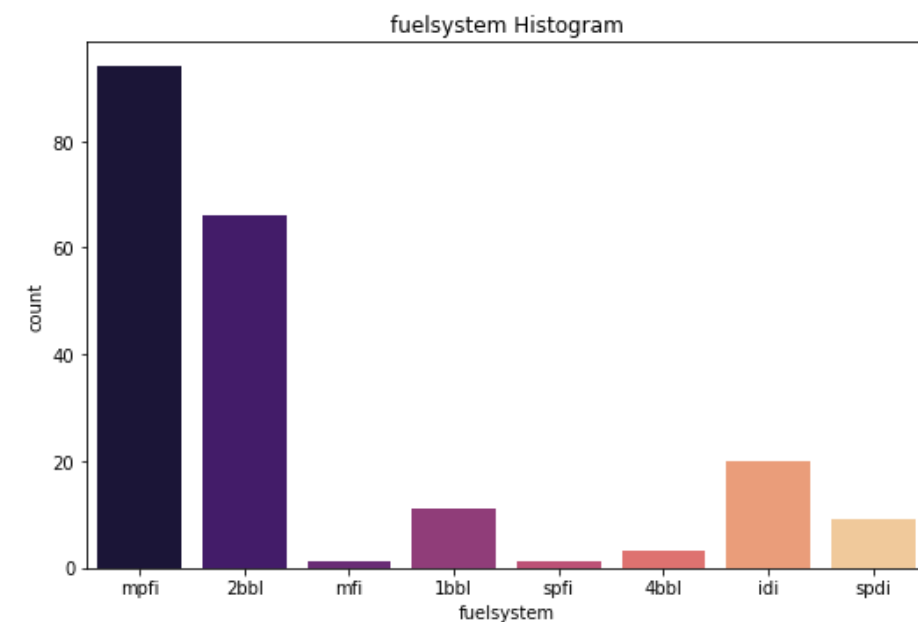
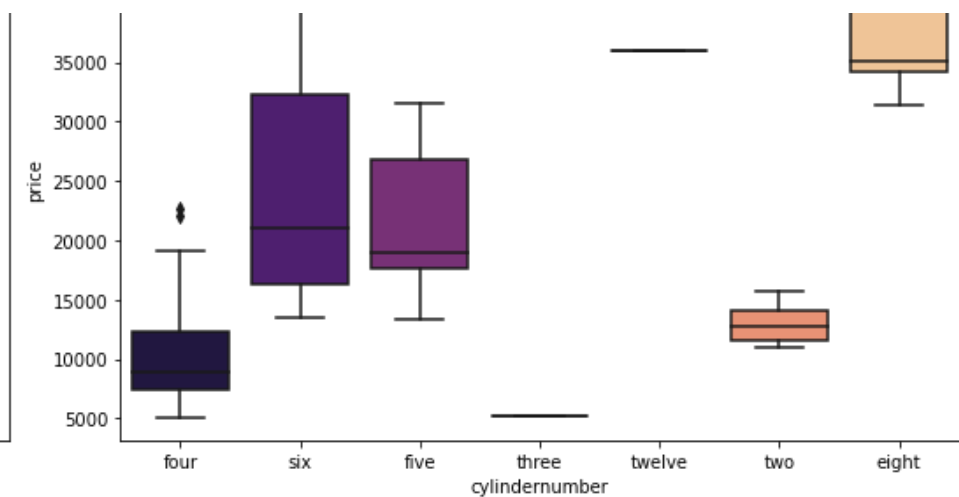
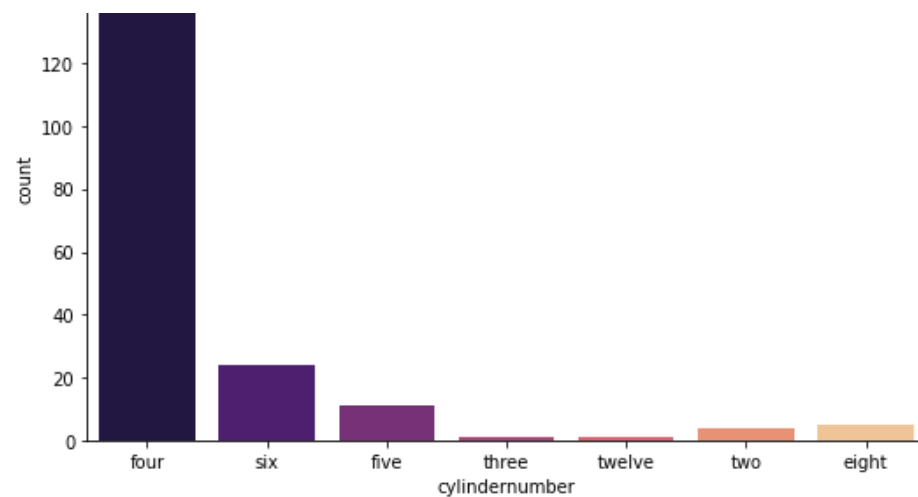
```
In [49]: def plot_count(x,fig):
plt.subplot(4,2,fig)
plt.title(x+' Histogram')
sns.countplot(cars[x],palette=("magma"))
plt.subplot(4,2,(fig+1))
plt.title(x+' vs Price')
sns.boxplot(x=cars[x], y=cars.price, palette=("magma"))


plt.figure(figsize=(15,20))

plot_count('enginelocation', 1)
plot_count('cylindernumber', 3)
plot_count('fuelsystem', 5)
plot_count('drivewheel', 7)

plt.tight_layout()
```







Inference :

1. Very few datapoints for engine location categories to make an inference.
2. Most common number of cylinders are four , six and five . Though eight cylinders have the highest price range.
3. mpfi and 2bbl are most common type of fuel systems. mpfi and idi having the highest price range. But there are few data for other categories to derive any meaningful inference
4. A very significant difference in drivewheel category. Most high ranged cars seem to prefer rwd drivewheel.

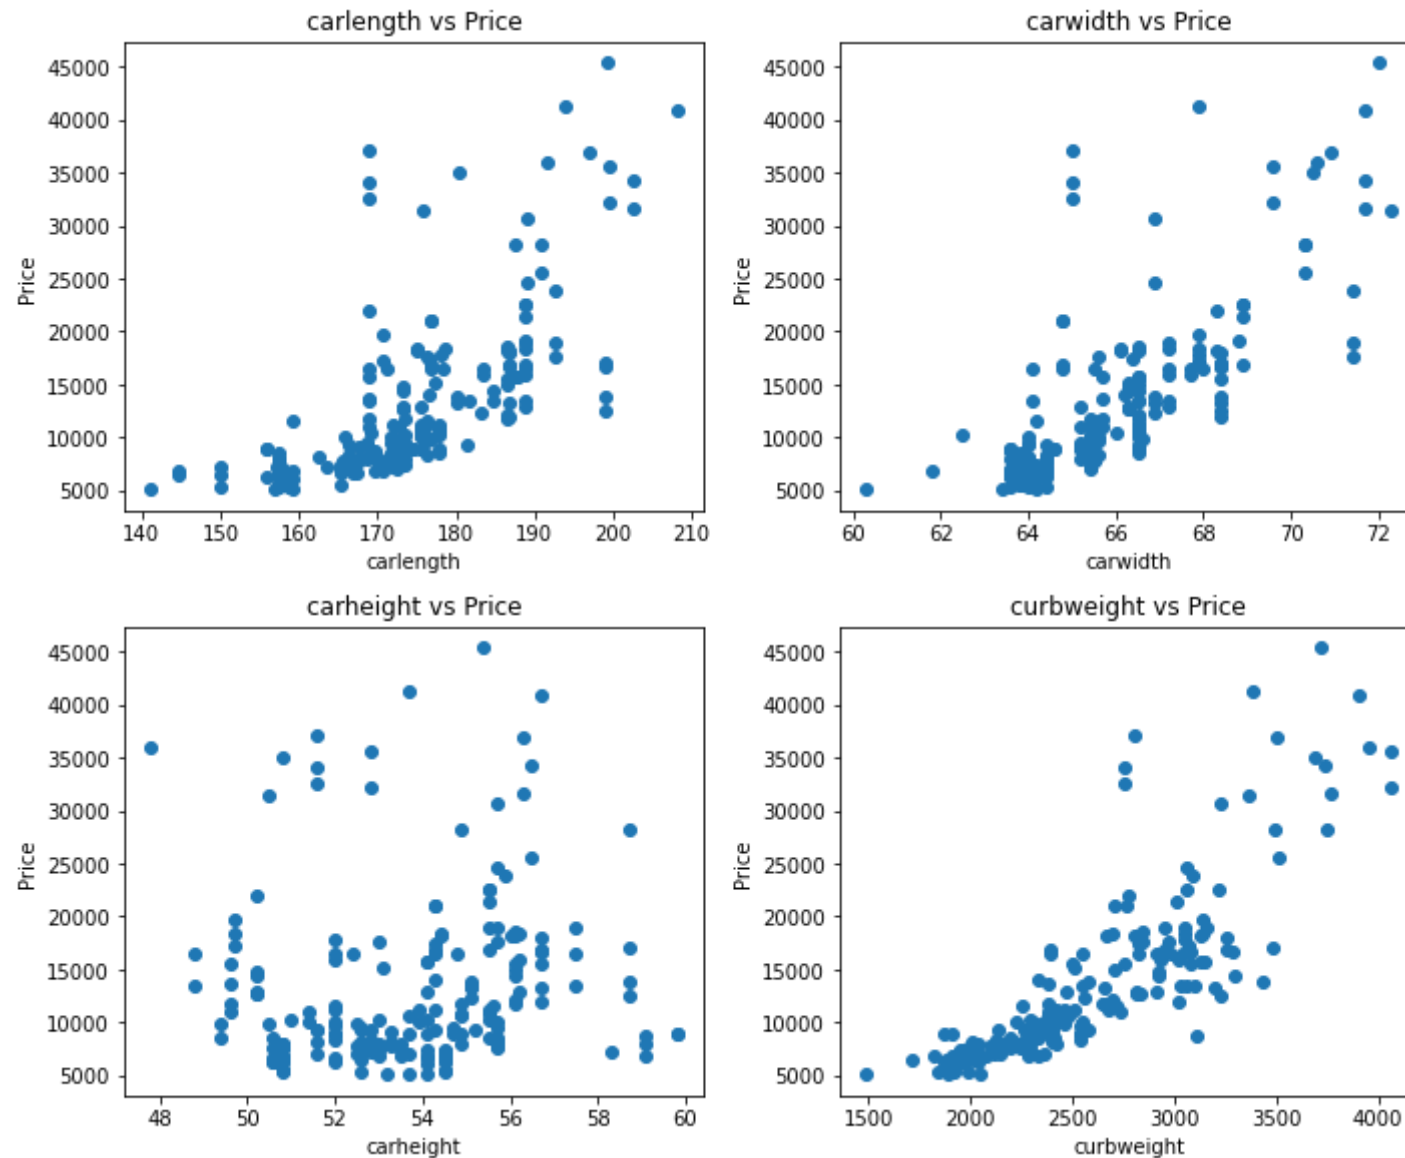
Step 2 : Visualising numerical data

```
In [50]: def scatter(x,fig):
    plt.subplot(5,2,fig)
    plt.scatter(cars[x],cars['price'])
    plt.title(x+' vs Price')
    plt.ylabel('Price')
    plt.xlabel(x)

plt.figure(figsize=(10,20))

scatter('carlength', 1)
scatter('carwidth', 2)
scatter('carheight', 3)
scatter('curbweight', 4)

plt.tight_layout()
```




Inference :


1. carwidth , carlength and curbweight seems to have a poitive correlation with price .
2. carheight doesn't show any significant trend with price.

```
In [51]: #Outliers also can be treated in this section
```

c. Do you want to exclude some variables from the model based on this analysis? What other actions will you take? (2 marks)

```
In [52]: def pp(x,y,z):  
    sns.pairplot(cars, x_vars=[x,y,z], y_vars='price',size=4, aspect=1, kind='scatter')  
    plt.show()  
  
pp('engine size', 'bore ratio', 'stroke')  
pp('compression ratio', 'horsepower', 'peak rpm')  
pp('wheelbase', 'city mpg', 'highway mpg')
```





Inference :

1. enginesize , boreratio , horsepower , wheelbase - seem to have a significant positive correlation with price.
2. citympg , highwaympg - seem to have a significant negative correlation with price.

```
In [ ]: np.corrcoef(cars['carlength'], cars['carwidth'])[0, 1]
```

Out[]: 0.841118268481846

features engineering

```
In [ ]: #Fuel economy
cars['fueleconomy'] = (0.55 * cars['citympg']) + (0.45 * cars['highwaympg'])
```

```
In [ ]: #Binning the Car Companies based on avg prices of each Company.
cars['price'] = cars['price'].astype('int')
temp = cars.copy()
table = temp.groupby(['CompanyName'])['price'].mean()
temp = temp.merge(table.reset_index(), how='left', on='CompanyName')
bins = [0,10000,20000,40000]
cars_bin=['Budget','Medium','Highend']
cars['carsrange'] = pd.cut(temp['price_y'],bins,right=False,labels=cars_bin)
cars.head()
```

```
Out[ ]:   car_ID  symboling  CompanyName  fueltype  aspiration  doornumber  carbody  drivewheel  enginelocation  wheelbase  ...  boreratio  stroke  compr
```

| | car_ID | symboling | CompanyName | fueltype | aspiration | doornumber | carbody | drivewheel | enginelocation | wheelbase | ... | boreratio | stroke | compr |
|---|--------|-----------|-------------|----------|------------|------------|-------------|------------|----------------|-----------|-----|-----------|--------|-------|
| 0 | 1 | 3 | alfa-romero | gas | std | two | convertible | rwd | front | 88.6 | ... | 3.47 | 2.68 | |
| 1 | 2 | 3 | alfa-romero | gas | std | two | convertible | rwd | front | 88.6 | ... | 3.47 | 2.68 | |
| 2 | 3 | 1 | alfa-romero | gas | std | two | hatchback | rwd | front | 94.5 | ... | 2.68 | 3.47 | |
| 3 | 4 | 2 | audi | gas | std | four | sedan | fwd | front | 99.8 | ... | 3.19 | 3.40 | |
| 4 | 5 | 2 | audi | gas | std | four | sedan | 4wd | front | 99.4 | ... | 3.19 | 3.40 | |

5 rows × 28 columns




Bivariate Analysis

```
In [ ]: plt.figure(figsize=(8,6))

plt.title('Fuel economy vs Price')
sns.scatterplot(x=cars['fueleconomy'],y=cars['price'],hue=cars['drivewheel'])
plt.xlabel('Fuel Economy')
plt.ylabel('Price')
```

```
plt.show()
plt.tight_layout()
```



<Figure size 432x288 with 0 Axes>


Inference :

1. fueleconomy has an obvious negative correlation with price and is significant.

```
In [ ]: plt.figure(figsize=(25, 6))

df = pd.DataFrame(cars.groupby(['fuelsystem', 'drivewheel', 'carsrange'])['price'].mean().unstack(fill_value=0))
df.plot.bar()
plt.title('Car Range vs Average Price')
plt.show()
```

<Figure size 1800x432 with 0 Axes>



Inference :

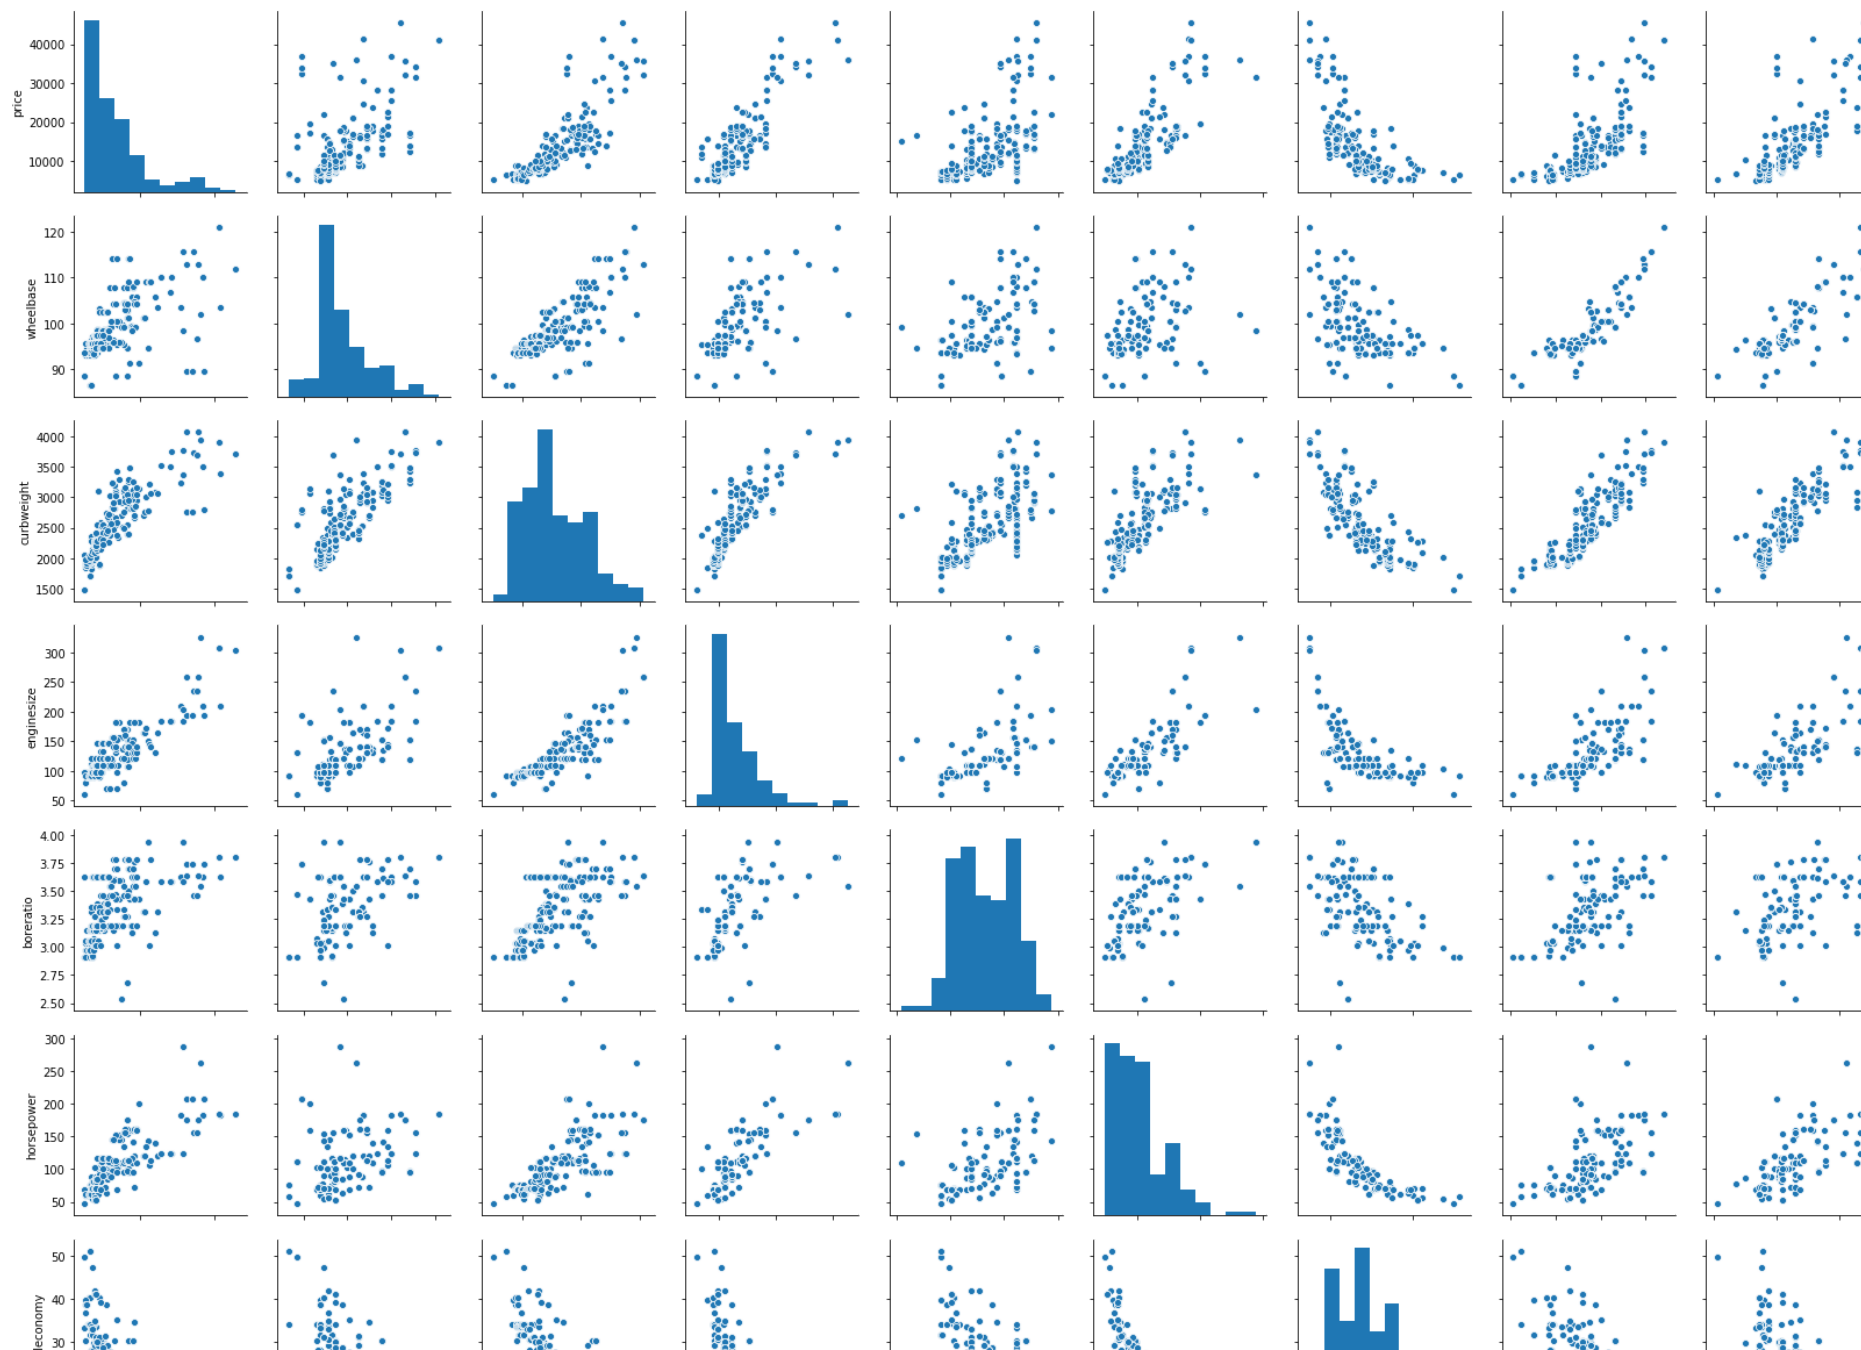
1. High ranged cars prefer rwd drivewheel with idi or mpfi fuelsystem.


```
In [ ]: cars_lr = cars[['price', 'fueltype', 'aspiration', 'carbody', 'drivewheel', 'wheelbase',
                        'curbweight', 'enginetype', 'cylindernumber', 'enginesize', 'boreratio', 'horsepower',
                        'fuel economy', 'carlength', 'carwidth', 'carsrange']]
cars_lr.head()
```

```
Out [ ]:
```

| | price | fueltype | aspiration | carbody | drivewheel | wheelbase | curbweight | enginetype | cylindernumber | enginesize | boreratio | horsepower | fuelcono |
|---|-------|----------|------------|-------------|------------|-----------|------------|------------|----------------|------------|-----------|------------|----------|
| 0 | 13495 | gas | std | convertible | rwd | 88.6 | 2548 | dohc | four | 130 | 3.47 | 111 | 23 |
| 1 | 16500 | gas | std | convertible | rwd | 88.6 | 2548 | dohc | four | 130 | 3.47 | 111 | 23 |
| 2 | 16500 | gas | std | hatchback | rwd | 94.5 | 2823 | ohcv | six | 152 | 2.68 | 154 | 22 |
| 3 | 13950 | gas | std | sedan | fwd | 99.8 | 2337 | ohc | four | 109 | 3.19 | 102 | 26 |
| 4 | 17450 | gas | std | sedan | 4wd | 99.4 | 2824 | ohc | five | 136 | 3.19 | 115 | 19 |

```
In [ ]: sns.pairplot(cars_lr)  
plt.show()
```





```
In [ ]: # Defining the map function
def dummies(x,df):
    temp = pd.get_dummies(df[x], drop_first = True)
    df = pd.concat([df, temp], axis = 1)
    df.drop([x], axis = 1, inplace = True)
    return df
# Applying the function to the cars_lr

cars_lr = dummies('fueltype',cars_lr)
cars_lr = dummies('aspiration',cars_lr)
cars_lr = dummies('carbody',cars_lr)
cars_lr = dummies('drivewheel',cars_lr)
cars_lr = dummies('enginetype',cars_lr)
cars_lr = dummies('cylindernumber',cars_lr)
cars_lr = dummies('carsrange',cars_lr)
```

```
In [ ]: cars_lr.head()
```

```
Out[ ]:
```

| | price | wheelbase | curbweight | enginesize | boreratio | horsepower | fuel economy | carlength | carwidth | gas | ... | ohcv | rotor | five | four | six | three | two |
|---|-------|-----------|------------|------------|-----------|------------|--------------|-----------|----------|-----|-----|------|-------|------|------|-----|-------|-----|
| 0 | 13495 | 88.6 | 2548 | 130 | 3.47 | 111 | 23.70 | 168.8 | 64.1 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 | |
| 1 | 16500 | 88.6 | 2548 | 130 | 3.47 | 111 | 23.70 | 168.8 | 64.1 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 | |
| 2 | 16500 | 94.5 | 2823 | 152 | 2.68 | 154 | 22.15 | 171.2 | 65.5 | 1 | ... | 1 | 0 | 0 | 0 | 1 | 0 | |
| 3 | 13950 | 99.8 | 2337 | 109 | 3.19 | 102 | 26.70 | 176.6 | 66.2 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 | |

| | price | wheelbase | curbweight | enginesize | boreratio | horsepower | fueleconomy | carlength | carwidth | gas | ... | ohcv | rotor | five | four | six | three | twe |
|---|-------|-----------|------------|------------|-----------|------------|-------------|-----------|----------|-----|-----|------|-------|------|------|-----|-------|-----|
| 4 | 17450 | 99.4 | 2824 | 136 | 3.19 | 115 | 19.80 | 176.6 | 66.4 | 1 | ... | 0 | 0 | 1 | 0 | 0 | 0 | |

5 rows × 31 columns



In []: `cars_lr.shape`

Out[]: (205, 31)

d. Split dataset into train and test (70:30). Are both train and test representative of the overall data? How would you ascertain this statistically? (3 marks)

```
In [ ]: from sklearn.model_selection import train_test_split

np.random.seed(0)
df_train, df_test = train_test_split(cars_lr, train_size = 0.7, test_size = 0.3, random_state = 100)
```

```
In [ ]: stats.ttest_ind(df_train.iloc[:,1:], df_test.iloc[:,1:])
#All the pvalues> 0.05
```

```
Out[ ]: Ttest_indResult(statistic=array([-0.84259363, -0.66959192, -1.01793949, -1.80371206, -0.85980387,
      0.4894942 , -1.11542253, -0.68781997,  0.48536645,  0.07485128,
      -2.03683  ,  0.0544803 ,  0.31367245,  0.72274638,  0.3971903 ,
      -0.94651433,  0.65754088,  0.40577066, -1.78317972,  0.89459542,
      -0.04240254,  1.32919754,  0.89270227, -1.79580298,  1.06593877,
      0.65754088,  0.65754088,  1.32919754, -0.30370793,  0.12168329]), pvalue=array([0.40044739, 0.50387885, 0.30991836, 0.0727
5905, 0.39091141,
      0.62502037, 0.26598809, 0.49235132, 0.62793983, 0.94040675,
      0.04296531, 0.95660608, 0.75409191, 0.47066756, 0.69164443,
      0.34501162, 0.51157807, 0.68533848, 0.0760512 , 0.37206285,
      0.96621948, 0.18527472, 0.37307364, 0.07401301, 0.28771745,
      0.51157807, 0.51157807, 0.18527472, 0.76166147, 0.90327021]))
```

```
In [ ]: stats.ttest_ind(df_train.iloc[:,0], df_test.iloc[:,0])
#All the pvalues> 0.05
```

```
Out[ ]: Ttest_indResult(statistic=-0.5988626606978698, pvalue=0.5499321581575893)
```

```
In [ ]: from sklearn.preprocessing import MinMaxScaler
```

```

scaler = MinMaxScaler()
num_vars = ['wheelbase', 'curbweight', 'enginesize', 'boreratio', 'horsepower', 'fuelconomy', 'carlength', 'carwidth', 'price']
df_train[num_vars] = scaler.fit_transform(df_train[num_vars])

```

In []: df_train.head()

Out []:

| | price | wheelbase | curbweight | enginesize | boreratio | horsepower | fuelconomy | carlength | carwidth | gas | ... | ohcv | rotor | five | four | six | three |
|-----|----------|-----------|------------|------------|-----------|------------|------------|-----------|----------|-----|-----|------|-------|------|------|-----|-------|
| 122 | 0.068818 | 0.244828 | 0.272692 | 0.139623 | 0.230159 | 0.083333 | 0.530864 | 0.426016 | 0.291667 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 |
| 125 | 0.466890 | 0.272414 | 0.500388 | 0.339623 | 1.000000 | 0.395833 | 0.213992 | 0.452033 | 0.666667 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 |
| 166 | 0.122110 | 0.272414 | 0.314973 | 0.139623 | 0.444444 | 0.266667 | 0.344307 | 0.448780 | 0.308333 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 |
| 1 | 0.314446 | 0.068966 | 0.411171 | 0.260377 | 0.626984 | 0.262500 | 0.244170 | 0.450407 | 0.316667 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 |
| 199 | 0.382131 | 0.610345 | 0.647401 | 0.260377 | 0.746032 | 0.475000 | 0.122085 | 0.775610 | 0.575000 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 |

5 rows × 31 columns



In []: df_test.head()

Out []:

| | price | wheelbase | curbweight | enginesize | boreratio | horsepower | fuelconomy | carlength | carwidth | gas | ... | ohcv | rotor | five | four | six | three | tw |
|-----|-------|-----------|------------|------------|-----------|------------|------------|-----------|----------|-----|-----|------|-------|------|------|-----|-------|----|
| 160 | 7738 | 95.7 | 2094 | 98 | 3.19 | 70 | 42.05 | 166.3 | 64.4 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 | |
| 186 | 8495 | 97.3 | 2275 | 109 | 3.19 | 85 | 30.15 | 171.7 | 65.5 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 | |
| 59 | 8845 | 98.8 | 2385 | 122 | 3.39 | 84 | 28.70 | 177.8 | 66.5 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 | |
| 165 | 9298 | 94.5 | 2265 | 98 | 3.24 | 112 | 27.35 | 168.7 | 64.0 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 | |
| 140 | 7603 | 93.3 | 2240 | 108 | 3.62 | 73 | 28.25 | 157.3 | 63.8 | 1 | ... | 0 | 0 | 0 | 1 | 0 | 0 | |

5 rows × 31 columns



In []: df_train.describe()

Out []:


| | price | wheelbase | curbweight | enginesize | boreratio | horsepower | fuelconomy | carlength | carwidth | gas | ... | ohcv | ro |
|--|-------|-----------|------------|------------|-----------|------------|------------|-----------|----------|-----|-----|------|----|
|--|-------|-----------|------------|------------|-----------|------------|------------|-----------|----------|-----|-----|------|----|


| | price | wheelbase | curbweight | enginesize | boreratio | horsepower | fueleconomy | carlength | carwidth | gas | ... | ohcv | ro |
|-------|------------|------------|------------|------------|------------|------------|-------------|------------|------------|------------|-----|------------|------------|
| count | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 | ... | 143.000000 | 143.000000 |
| mean | 0.219309 | 0.411141 | 0.407878 | 0.241351 | 0.497946 | 0.227302 | 0.358265 | 0.525476 | 0.461655 | 0.909091 | ... | 0.062937 | 0.027937 |
| std | 0.215682 | 0.205581 | 0.211269 | 0.154619 | 0.207140 | 0.165511 | 0.185980 | 0.204848 | 0.184517 | 0.288490 | ... | 0.243703 | 0.165403 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | ... | 0.000000 | 0.000000 |
| 25% | 0.067298 | 0.272414 | 0.245539 | 0.135849 | 0.305556 | 0.091667 | 0.198903 | 0.399187 | 0.304167 | 1.000000 | ... | 0.000000 | 0.000000 |
| 50% | 0.140343 | 0.341379 | 0.355702 | 0.184906 | 0.500000 | 0.191667 | 0.344307 | 0.502439 | 0.425000 | 1.000000 | ... | 0.000000 | 0.000000 |
| 75% | 0.313479 | 0.503448 | 0.559542 | 0.301887 | 0.682540 | 0.283333 | 0.512346 | 0.669919 | 0.550000 | 1.000000 | ... | 0.000000 | 0.000000 |
| max | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | ... | 1.000000 | 1.000000 |

8 rows × 31 columns



```
In [ ]: #Correlation using heatmap
plt.figure(figsize = (30, 25))
sns.heatmap(df_train.corr(), annot = True, cmap="YlGnBu")
plt.show()
```





Highly correlated variables to price are - curbweight , enginesize , horsepower , carwidth and highend .

```
In [ ]: #Dividing data into X and y variables
y_train = df_train.pop('price')
X_train = df_train
```

```
In [ ]: from scipy import stats
```

3. Model Building (20 marks)

- Fit a base model and observe the overall R- Squared, RMSE and MAPE values of the model. Please comment on whether it is good or not. (5 marks)
- Check for multi-collinearity and treat the same. (3 marks)
- How would you improve the model? Write clearly the changes that you will make before re-fitting the model. Fit the final model. (6 marks)
- Write down a business interpretation/explanation of the model – which variables are affecting the target the most and explain the relationship. Feel free to use charts or graphs to explain. (4 marks)
- What changes from the base model had the most effect on model performance? (2 marks)

a. Fit a base model and observe the overall R- Squared, RMSE and MAPE values of the model. Please comment on whether it is good or not. (5 marks)

```
In [ ]: #RFE
        from sklearn.feature_selection import RFE
        from sklearn.linear_model import LinearRegression
        import statsmodels.api as sm
        from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [ ]: lm = LinearRegression()
        lm.fit(X_train,y_train)
        rfe = RFE(lm, 10)
        rfe = rfe.fit(X_train, y_train)
```

```
In [ ]: X_train.columns[rfe.support_]
```

```
Out[ ]: Index(['curbweight', 'horsepower', 'fueleconomy', 'carwidth', 'hatchback',
              'sedan', 'wagon', 'dohcv', 'twelve', 'Highend'],
              dtype='object')
```

Building model using statsmodel, for the detailed statistics

```
In [ ]: X_train_rfe = X_train[X_train.columns[rfe.support_]]
        X_train_rfe.head()
```

```
Out[ ]:      curbweight  horsepower  fueleconomy  carwidth  hatchback  sedan  wagon  dohcv  twelve  Highend
122      0.272692      0.083333      0.530864  0.291667           0       1       0       0       0       0
```


| | curbweight | horsepower | fuel economy | carwidth | hatchback | sedan | wagon | dohcv | twelve | Highend |
|------------|------------|------------|--------------|----------|-----------|-------|-------|-------|--------|---------|
| 125 | 0.500388 | 0.395833 | 0.213992 | 0.666667 | 1 | 0 | 0 | 0 | 0 | 1 |
| 166 | 0.314973 | 0.266667 | 0.344307 | 0.308333 | 1 | 0 | 0 | 0 | 0 | 0 |
| 1 | 0.411171 | 0.262500 | 0.244170 | 0.316667 | 0 | 0 | 0 | 0 | 0 | 0 |
| 199 | 0.647401 | 0.475000 | 0.122085 | 0.575000 | 0 | 0 | 1 | 0 | 0 | 0 |

```
In [ ]: def build_model(X,y):
        X = sm.add_constant(X) #Adding the constant
        lm = sm.OLS(y,X).fit() # fitting the model
        print(lm.summary()) # model summary
        return X

def checkVIF(X):
    vif = pd.DataFrame()
    vif['Features'] = X.columns
    vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
    vif['VIF'] = round(vif['VIF'], 2)
    vif = vif.sort_values(by = "VIF", ascending = False)
    return(vif)
```

MODEL 1

```
In [ ]: X_train_new = build_model(X_train_rfe,y_train)
```

```

OLS Regression Results
=====
Dep. Variable:          price    R-squared:                0.929
Model:                  OLS      Adj. R-squared:           0.923
Method:                 Least Squares    F-statistic:          172.1
Date:                  Wed, 22 Jul 2020    Prob (F-statistic):    1.29e-70
Time:                  13:58:33      Log-Likelihood:        205.85
No. Observations:      143          AIC:                   -389.7
Df Residuals:          132          BIC:                   -357.1
Df Model:              10
Covariance Type:       nonrobust
=====
               coef    std err          t      P>|t|      [0.025    0.975]
-----
const         -0.0947     0.042     -2.243     0.027     -0.178     -0.011
curbweight      0.2657     0.069      3.870     0.000      0.130      0.402

```

| | | | | | | |
|------------|---------|-------|--------|-------|--------|--------|
| horsepower | 0.4499 | 0.074 | 6.099 | 0.000 | 0.304 | 0.596 |
| fuelconomy | 0.0933 | 0.052 | 1.792 | 0.075 | -0.010 | 0.196 |
| carwidth | 0.2609 | 0.062 | 4.216 | 0.000 | 0.138 | 0.383 |
| hatchback | -0.0929 | 0.025 | -3.707 | 0.000 | -0.143 | -0.043 |
| sedan | -0.0704 | 0.025 | -2.833 | 0.005 | -0.120 | -0.021 |
| wagon | -0.0997 | 0.028 | -3.565 | 0.001 | -0.155 | -0.044 |
| dohcv | -0.2676 | 0.079 | -3.391 | 0.001 | -0.424 | -0.112 |
| twelve | -0.1192 | 0.067 | -1.769 | 0.079 | -0.253 | 0.014 |
| Highend | 0.2586 | 0.020 | 12.929 | 0.000 | 0.219 | 0.298 |

| | | | |
|----------------|--------|-------------------|----------|
| Omnibus: | 43.093 | Durbin-Watson: | 1.867 |
| Prob(Omnibus): | 0.000 | Jarque-Bera (JB): | 130.648 |
| Skew: | 1.128 | Prob(JB): | 4.27e-29 |
| Kurtosis: | 7.103 | Cond. No. | 32.0 |

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

p-vale of twelve seems to be higher than the significance value of 0.05, hence dropping it as it is insignificant in presence of other variables.

Here overall R- Squared

- R-squared: 0.929 Adj. R-squared: 0.923

```
In [ ]: X_train_new = X_train_rfe.drop(["twelve"], axis = 1)
```

```
In [ ]: X_train_new = build_model(X_train_new,y_train)
```

OLS Regression Results

| | | | |
|-------------------|------------------|---------------------|----------|
| Dep. Variable: | price | R-squared: | 0.927 |
| Model: | OLS | Adj. R-squared: | 0.922 |
| Method: | Least Squares | F-statistic: | 187.9 |
| Date: | Wed, 22 Jul 2020 | Prob (F-statistic): | 4.25e-71 |
| Time: | 13:58:33 | Log-Likelihood: | 204.17 |
| No. Observations: | 143 | AIC: | -388.3 |
| Df Residuals: | 133 | BIC: | -358.7 |
| Df Model: | 9 | | |
| Covariance Type: | nonrobust | | |

| | coef | std err | t | P> t | [0.025 | 0.975] |
|------------|---------|---------|--------|-------|--------|--------|
| const | -0.0764 | 0.041 | -1.851 | 0.066 | -0.158 | 0.005 |
| curbweight | 0.2756 | 0.069 | 3.995 | 0.000 | 0.139 | 0.412 |
| horsepower | 0.3997 | 0.069 | 5.824 | 0.000 | 0.264 | 0.535 |

| | | | | | | |
|-------------|---------|-------|--------|-------|--------|--------|
| fueleconomy | 0.0736 | 0.051 | 1.435 | 0.154 | -0.028 | 0.175 |
| carwidth | 0.2580 | 0.062 | 4.137 | 0.000 | 0.135 | 0.381 |
| hatchback | -0.0951 | 0.025 | -3.766 | 0.000 | -0.145 | -0.045 |
| sedan | -0.0744 | 0.025 | -2.983 | 0.003 | -0.124 | -0.025 |
| wagon | -0.1050 | 0.028 | -3.744 | 0.000 | -0.160 | -0.050 |
| dohcv | -0.2319 | 0.077 | -3.015 | 0.003 | -0.384 | -0.080 |
| Highend | 0.2565 | 0.020 | 12.743 | 0.000 | 0.217 | 0.296 |

```
=====
Omnibus:                48.027    Durbin-Watson:                1.880
Prob(Omnibus):          0.000    Jarque-Bera (JB):        159.802
Skew:                   1.231    Prob(JB):                1.99e-35
Kurtosis:                7.556    Cond. No.                29.6
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [ ]: X_train_new = X_train_new.drop(["fueleconomy"], axis = 1) #pvalue>0.05
```

MODEL 3

```
In [ ]: X_train_new = build_model(X_train_new,y_train)
```

OLS Regression Results

```
=====
Dep. Variable:          price    R-squared:                0.926
Model:                  OLS     Adj. R-squared:            0.922
Method:                 Least Squares    F-statistic:            209.5
Date:                   Wed, 22 Jul 2020    Prob (F-statistic):      7.85e-72
Time:                   13:58:34    Log-Likelihood:          203.07
No. Observations:        143    AIC:                    -388.1
Df Residuals:            134    BIC:                    -361.5
Df Model:                 8
Covariance Type:         nonrobust
=====
```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|------------|---------|---------|--------|-------|--------|--------|
| const | -0.0305 | 0.026 | -1.165 | 0.246 | -0.082 | 0.021 |
| curbweight | 0.2593 | 0.068 | 3.796 | 0.000 | 0.124 | 0.394 |
| horsepower | 0.3469 | 0.058 | 5.964 | 0.000 | 0.232 | 0.462 |
| carwidth | 0.2488 | 0.062 | 3.995 | 0.000 | 0.126 | 0.372 |
| hatchback | -0.0922 | 0.025 | -3.650 | 0.000 | -0.142 | -0.042 |
| sedan | -0.0711 | 0.025 | -2.850 | 0.005 | -0.120 | -0.022 |
| wagon | -0.1047 | 0.028 | -3.721 | 0.000 | -0.160 | -0.049 |
| dohcv | -0.1968 | 0.073 | -2.689 | 0.008 | -0.342 | -0.052 |
| Highend | 0.2610 | 0.020 | 13.083 | 0.000 | 0.222 | 0.301 |

```
=====
Omnibus:                48.637    Durbin-Watson:                1.909
Prob(Omnibus):           0.000    Jarque-Bera (JB):            161.444
Skew:                    1.250    Prob(JB):                     8.77e-36
Kurtosis:                7.566    Cond. No.                     27.2
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

variables are significant?

List of significant variables :

- Car Range
- Engine Type
- Fuel type
- Car Body
- Aspiration
- Cylinder Number
- Drivewheel
- Curbweight
- Car Length
- Car width
- Engine Size
- Boreratio
- Horse Power
- Wheel base
- Fuel Economy

b. Check for multi-collinearity and treat the same. (3 marks)

```
In [ ]: #Calculating the Variance Inflation Factor
        checkVIF(X_train_new)
```

```
Out[ ]:
```

| | Features | VIF |
|---|------------|-------|
| 0 | const | 26.90 |
| 1 | curbweight | 8.10 |
| 5 | sedan | 6.07 |

| | Features | VIF |
|---|------------|------|
| 4 | hatchback | 5.63 |
| 3 | carwidth | 5.14 |
| 2 | horsepower | 3.61 |
| 6 | wagon | 3.58 |
| 8 | Highend | 1.63 |
| 7 | dohcv | 1.46 |

dropping `curbweight` because of high VIF value. (shows that curbweight has high multicollinearity.)

```
In [ ]: X_train_new = X_train_new.drop(["curbweight"], axis = 1)
```

c. How would you improve the model? Write clearly the changes that you will make before re-fitting the model. Fit the final model. (6 marks)

```
In [ ]: X_train_new = build_model(X_train_new,y_train)
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.918
Model:                  OLS       Adj. R-squared:            0.914
Method:                 Least Squares   F-statistic:            215.9
Date:                  Wed, 22 Jul 2020   Prob (F-statistic):      4.70e-70
Time:                  13:58:37          Log-Likelihood:         195.77
No. Observations:      143             AIC:                   -375.5
Df Residuals:          135             BIC:                   -351.8
Df Model:               7
Covariance Type:       nonrobust
=====

```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|------------|---------|---------|--------|-------|--------|--------|
| const | -0.0319 | 0.027 | -1.161 | 0.248 | -0.086 | 0.022 |
| horsepower | 0.4690 | 0.051 | 9.228 | 0.000 | 0.368 | 0.569 |
| carwidth | 0.4269 | 0.043 | 9.944 | 0.000 | 0.342 | 0.512 |
| hatchback | -0.1044 | 0.026 | -3.976 | 0.000 | -0.156 | -0.052 |
| sedan | -0.0756 | 0.026 | -2.896 | 0.004 | -0.127 | -0.024 |
| wagon | -0.0865 | 0.029 | -2.974 | 0.003 | -0.144 | -0.029 |
| dohcv | -0.3106 | 0.070 | -4.435 | 0.000 | -0.449 | -0.172 |
| Highend | 0.2772 | 0.020 | 13.559 | 0.000 | 0.237 | 0.318 |

```
=====
Omnibus:                43.937    Durbin-Watson:                2.006
Prob(Omnibus):           0.000    Jarque-Bera (JB):          127.746
Skew:                    1.171    Prob(JB):                  1.82e-28
Kurtosis:                6.995    Cond. No.                  18.0
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In []: `checkVIF(X_train_new)`

Out []:

| | Features | VIF |
|---|------------|-------|
| 0 | const | 26.89 |
| 4 | sedan | 6.06 |
| 3 | hatchback | 5.54 |
| 5 | wagon | 3.47 |
| 1 | horsepower | 2.50 |
| 2 | carwidth | 2.22 |
| 7 | Highend | 1.56 |
| 6 | dohcv | 1.21 |

dropping sedan because of high VIF value.

In []: `X_train_new = X_train_new.drop(["sedan"], axis = 1)`

MODEL

In []: `X_train_new = build_model(X_train_new,y_train)`

```

                        OLS Regression Results
=====
Dep. Variable:          price    R-squared:                0.913
Model:                  OLS      Adj. R-squared:            0.909
Method:                 Least Squares    F-statistic:           237.6
Date:                  Wed, 22 Jul 2020    Prob (F-statistic):     1.68e-69
Time:                  13:58:38    Log-Likelihood:        191.46
No. Observations:      143    AIC:                   -368.9

```

```

Df Residuals:      136    BIC:      -348.2
Df Model:           6
Covariance Type:    nonrobust

```

```

=====
              coef    std err          t      P>|t|      [0.025      0.975]
-----
const         -0.0934     0.018    -5.219     0.000    -0.129    -0.058
horsepower      0.5001     0.051     9.805     0.000     0.399     0.601
carwidth        0.3963     0.043     9.275     0.000     0.312     0.481
hatchback      -0.0373     0.013    -2.938     0.004    -0.062    -0.012
wagon          -0.0170     0.017    -1.008     0.315    -0.050     0.016
dohcv          -0.3203     0.072    -4.460     0.000    -0.462    -0.178
Highend         0.2808     0.021    13.402     0.000     0.239     0.322
=====

```

```

Omnibus:      34.143    Durbin-Watson:      2.024
Prob(Omnibus): 0.000    Jarque-Bera (JB):      72.788
Skew:         1.018    Prob(JB):      1.56e-16
Kurtosis:     5.841    Cond. No.      16.4
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [ ]: checkVIF(X_train_new)
```

```

Out[ ]:

```

| | Features | VIF |
|---|------------|-------|
| 0 | const | 10.82 |
| 1 | horsepower | 2.39 |
| 2 | carwidth | 2.09 |
| 6 | Highend | 1.55 |
| 3 | hatchback | 1.23 |
| 5 | dohcv | 1.21 |
| 4 | wagon | 1.11 |

dropping wagon because of high p-value.

```
In [ ]: X_train_new = X_train_new.drop(["wagon"], axis = 1)
```

MODEL

```
In [ ]: X_train_new = build_model(X_train_new,y_train)
```

```

                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.912
Model:                  OLS        Adj. R-squared:            0.909
Method:                 Least Squares    F-statistic:           284.8
Date:                  Wed, 22 Jul 2020    Prob (F-statistic):    1.57e-70
Time:                  13:58:39      Log-Likelihood:        190.93
No. Observations:      143          AIC:                   -369.9
Df Residuals:          137          BIC:                   -352.1
Df Model:               5
Covariance Type:       nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const         -0.0970      0.018     -5.530      0.000     -0.132     -0.062
horsepower      0.5013      0.051      9.832      0.000      0.401      0.602
carwidth        0.3952      0.043      9.252      0.000      0.311      0.480
hatchback       -0.0336      0.012     -2.764      0.006     -0.058     -0.010
dohcv           -0.3231      0.072     -4.502      0.000     -0.465     -0.181
Highend         0.2833      0.021     13.615      0.000      0.242      0.324
=====
Omnibus:            36.097    Durbin-Watson:           2.028
Prob(Omnibus):      0.000    Jarque-Bera (JB):        78.717
Skew:               1.067    Prob(JB):                8.07e-18
Kurtosis:           5.943    Cond. No.                 16.3
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [ ]: checkVIF(X_train_new)
```

```
Out[ ]:
```

| | Features | VIF |
|---|------------|-------|
| 0 | const | 10.39 |
| 1 | horsepower | 2.39 |
| 2 | carwidth | 2.08 |
| 5 | Highend | 1.53 |
| 4 | dohcv | 1.21 |

| | Features | VIF |
|---|-----------|------|
| 3 | hatchback | 1.13 |

MODEL

```
In [ ]: #Dropping dohcvc to see the changes in model statistics
X_train_new = X_train_new.drop(["dohcvc"], axis = 1)
X_train_new = build_model(X_train_new,y_train)
checkVIF(X_train_new)
```

```
=====
                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.899
Model:                  OLS       Adj. R-squared:            0.896
Method:                 Least Squares   F-statistic:           308.0
Date:                  Wed, 22 Jul 2020   Prob (F-statistic):    1.04e-67
Time:                  13:58:40         Log-Likelihood:        181.06
No. Observations:      143             AIC:                  -352.1
Df Residuals:          138             BIC:                  -337.3
Df Model:              4
Covariance Type:       nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const          -0.0824      0.018      -4.480      0.000      -0.119      -0.046
horsepower      0.4402      0.052       8.390      0.000       0.336       0.544
carwidth        0.3957      0.046       8.677      0.000       0.306       0.486
hatchback      -0.0414      0.013      -3.219      0.002      -0.067      -0.016
Highend         0.2794      0.022      12.591      0.000       0.236       0.323
=====
Omnibus:            29.385   Durbin-Watson:           1.955
Prob(Omnibus):      0.000   Jarque-Bera (JB):        98.010
Skew:               0.692   Prob(JB):                5.22e-22
Kurtosis:           6.812   Cond. No.                 12.9
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
Out[ ]:
Features  VIF
0      const 10.04
1 horsepower 2.22
```


| | Features | VIF |
|---|-----------|------|
| 2 | carwidth | 2.08 |
| 4 | Highend | 1.53 |
| 3 | hatchback | 1.10 |

Residual Analysis of Model

```
In [ ]: lm = sm.OLS(y_train,X_train_new).fit()
        y_train_price = lm.predict(X_train_new)
```

```
In [ ]: # Plot the histogram of the error terms
        fig = plt.figure()
        sns.distplot((y_train - y_train_price), bins = 20)
        fig.suptitle('Error Terms', fontsize = 20)           # Plot heading
        plt.xlabel('Errors', fontsize = 18)
```

```
Out[ ]: Text(0.5, 0, 'Errors')
```



Error terms seem to be approximately normally distributed, so the assumption on the linear modeling seems to be fulfilled.

Prediction and Evaluation

```
In [ ]: #Scaling the test set
num_vars = ['wheelbase', 'curbweight', 'enginesize', 'boreatio', 'horsepower', 'fueleconomy', 'carlength', 'carwidth', 'price']
df_test[num_vars] = scaler.fit_transform(df_test[num_vars])
```

```
In [ ]: #Dividing into X and y
y_test = df_test.pop('price')
X_test = df_test
```

```
In [ ]: # Now let's use our model to make predictions.
X_train_new = X_train_new.drop('const',axis=1)
# Creating X_test_new dataframe by dropping variables from X_test
X_test_new = X_test[X_train_new.columns]

# Adding a constant variable
X_test_new = sm.add_constant(X_test_new)
```

```
In [ ]: # Making predictions
y_pred = lm.predict(X_test_new)
```


Evaluation of test via comparison of y_pred and y_test

```
In [ ]: from sklearn.metrics import r2_score
r2_score(y_test, y_pred)
```

```
Out[ ]: 0.8614595209022033
```

```
In [ ]: #EVALUATION OF THE MODEL
# Plotting y_test and y_pred to understand the spread.
fig = plt.figure()
plt.scatter(y_test,y_pred)
fig.suptitle('y_test vs y_pred', fontsize=20)           # Plot heading
plt.xlabel('y_test', fontsize=18)                      # X-Label
plt.ylabel('y_pred', fontsize=16)
```

```
Out[ ]: Text(0, 0.5, 'y_pred')
```



Evaluation of the model using Statistics

```
In [ ]: print(lm.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.899
Model:                  OLS       Adj. R-squared:            0.896
Method:                 Least Squares   F-statistic:             308.0
Date:                   Wed, 22 Jul 2020   Prob (F-statistic):      1.04e-67
Time:                   13:58:44    Log-Likelihood:          181.06
No. Observations:       143          AIC:                    -352.1
Df Residuals:           138          BIC:                    -337.3
Df Model:                4
Covariance Type:        nonrobust
=====

```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|------------|---------|---------|--------|-------|--------|--------|
| const | -0.0824 | 0.018 | -4.480 | 0.000 | -0.119 | -0.046 |
| horsepower | 0.4402 | 0.052 | 8.390 | 0.000 | 0.336 | 0.544 |
| carwidth | 0.3957 | 0.046 | 8.677 | 0.000 | 0.306 | 0.486 |
| hatchback | -0.0414 | 0.013 | -3.219 | 0.002 | -0.067 | -0.016 |
| Highend | 0.2794 | 0.022 | 12.591 | 0.000 | 0.236 | 0.323 |

```

=====
Omnibus:                 29.385    Durbin-Watson:           1.955
Prob(Omnibus):            0.000    Jarque-Bera (JB):        98.010

```

| | | | |
|-----------|-------|-----------|----------|
| Skew: | 0.692 | Prob(JB): | 5.22e-22 |
| Kurtosis: | 6.812 | Cond. No. | 12.9 |
| ===== | | | |

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

d. Write down a business interpretation/explanation of the model – which variables are affecting the target the most and explain the relationship. Feel free to use charts or graphs to explain. (4 marks)

1. *R-squared and Adjusted R-squared (extent of fit)* - 0.899 and 0.896 - 90% variance explained.
2. *F-stats and Prob(F-stats) (overall model fit)* - 308.0 and 1.04e-67(approx. 0.0) - Model fit is significant and explained 90% variance is just not by chance.
3. *p-values* - p-values for all the coefficients seem to be less than the significance level of 0.05. - meaning that all the predictors are statistically significant.

In []: